

# Simple Linear Regression

```
pip install numpy
```

Requirement already satisfied: numpy in c:\anaconda\lib\site-packages (1.21.5)

Note: you may need to restart the kernel to use updated packages.

```
pip install pandas
```

Requirement already satisfied: pandas in c:\anaconda\lib\site-packages (1.4.2) Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: python-dateutil>=2.8.1 in c:\anaconda\lib\site-packages (from pandas) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in c:\anaconda\lib\site-packages (from pandas) (2021.3)

Requirement already satisfied: numpy>=1.18.5 in c:\anaconda\lib\site-packages (from pandas) (1.21.5)

Requirement already satisfied: six>=1.5 in c:\anaconda\lib\site-packages (from python-dateutil>=2.8.1->pandas) (1.16.0)

```
pip install -U scikit-learn
```

Requirement already satisfied: scikit-learn in c:\anaconda\lib\site-packages (1.0.2)

Collecting scikit-learn

Downloading scikit\_learn-1.3.0-cp39-cp39-win\_amd64.whl (9.3 MB)

Requirement already satisfied: numpy>=1.17.3 in c:\anaconda\lib\site-packages (from scikit-learn) (1.21.5)

Requirement already satisfied: threadpoolctl>=2.0.0 in c:\anaconda\lib\site-packages (from scikit-learn) (2.2.0)

Collecting joblib>=1.1.1

Downloading joblib-1.3.1-py3-none-any.whl (301 kB)

Requirement already satisfied: scipy>=1.5.0 in c:\anaconda\lib\site-packages (from scikit-learn) (1.7.3)

Installing collected packages: joblib, scikit-learn

Attempting uninstall: joblib

Found existing installation: joblib 1.1.0

Uninstalling joblib-1.1.0:

Successfully uninstalled joblib-1.1.0

Attempting uninstall: scikit-learn

Found existing installation: scikit-learn 1.0.2

Uninstalling scikit-learn-1.0.2:

Successfully uninstalled scikit-learn-1.0.2

Successfully installed joblib-1.3.1 scikit-learn-1.3.0

Note: you may need to restart the kernel to use updated packages.

```
pip install matplotlib
```

```
Requirement already satisfied: matplotlib in c:\anaconda\lib\site-packages (3.5.1)
```

```
Requirement already satisfied: cyclor>=0.10 in c:\anaconda\lib\site-packages (from matplotlib) (0.11.0)
```

```
Requirement already satisfied: fonttools>=4.22.0 in c:\anaconda\lib\site-packages (from matplotlib) (4.25.0)
```

```
Requirement already satisfied: packaging>=20.0 in c:\anaconda\lib\site-packages (from matplotlib) (21.3)
```

```
Requirement already satisfied: numpy>=1.17 in c:\anaconda\lib\site-packages (from matplotlib) (1.21.5)
```

```
Requirement already satisfied: kiwisolver>=1.0.1 in c:\anaconda\lib\site-packages (from matplotlib) (1.3.2)
```

```
Requirement already satisfied: pyparsing>=2.2.1 in c:\anaconda\lib\site-packages (from matplotlib) (3.0.4)
```

```
Requirement already satisfied: pillow>=6.2.0 in c:\anaconda\lib\site-packages (from matplotlib) (9.0.1)
```

```
Requirement already satisfied: python-dateutil>=2.7 in c:\anaconda\lib\site-packages (from matplotlib) (2.8.2)
```

```
Requirement already satisfied: six>=1.5 in c:\anaconda\lib\site-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
```

```
Note: you may need to restart the kernel to use updated packages.
```

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
data = pd.read_csv(r'''C:\Users\Sanika\Downloads\Salary_Data.csv''')
```

```
data
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0

17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0
25	9.0	105582.0
26	9.5	116969.0
27	9.6	112635.0
28	10.3	122391.0
29	10.5	121872.0

```
data.head()
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

```
data.tail()
```

	YearsExperience	Salary
25	9.0	105582.0
26	9.5	116969.0
27	9.6	112635.0
28	10.3	122391.0
29	10.5	121872.0

```
data[13:21]
```

	YearsExperience	Salary
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0

```
x = data.iloc[:, :-1].values
y = data.iloc[:, 1:].values
```

```
x
```

```
array([[ 1.1],
       [ 1.3],
```

```
[ 1.5],  
[ 2. ],  
[ 2.2],  
[ 2.9],  
[ 3. ],  
[ 3.2],  
[ 3.2],  
[ 3.7],  
[ 3.9],  
[ 4. ],  
[ 4. ],  
[ 4.1],  
[ 4.5],  
[ 4.9],  
[ 5.1],  
[ 5.3],  
[ 5.9],  
[ 6. ],  
[ 6.8],  
[ 7.1],  
[ 7.9],  
[ 8.2],  
[ 8.7],  
[ 9. ],  
[ 9.5],  
[ 9.6],  
[10.3],  
[10.5]])
```

y

```
array([[ 39343.],  
[ 46205.],  
[ 37731.],  
[ 43525.],  
[ 39891.],  
[ 56642.],  
[ 60150.],  
[ 54445.],  
[ 64445.],  
[ 57189.],  
[ 63218.],  
[ 55794.],  
[ 56957.],  
[ 57081.],  
[ 61111.],  
[ 67938.],  
[ 66029.],  
[ 83088.],  
[ 81363.]])
```

```
[ 93940.],  
[ 91738.],  
[ 98273.],  
[101302.],  
[113812.],  
[109431.],  
[105582.],  
[116969.],  
[112635.],  
[122391.],  
[121872.]])
```

```
from sklearn.model_selection import train_test_split  
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size =  
1/3, random_state = 0)
```

```
from sklearn.linear_model import LinearRegression  
regressor = LinearRegression()  
regressor.fit(x_train, y_train)
```

```
LinearRegression()
```

```
y_pred = regressor.predict(x_test)
```

```
plt.scatter(x_train, y_train, color = 'green')  
plt.plot(x_train, regressor.predict(x_train), color = 'blue')  
plt.title("Salary vs Experiance (Training set)")  
plt.xlabel("Years of Experiance")  
plt.ylabel("Salary")  
plt.show()
```



```
plt.scatter(x_test, y_test, color = 'red')
plt.plot(x_train, regressor.predict(x_train), color = 'violet')
plt.title("Salary vs Experiance (Test set)")
plt.xlabel("Years of Experiance")
plt.ylabel("Salary")
plt.show()
```



# Multiple linear regression

```
import numpy as np
```

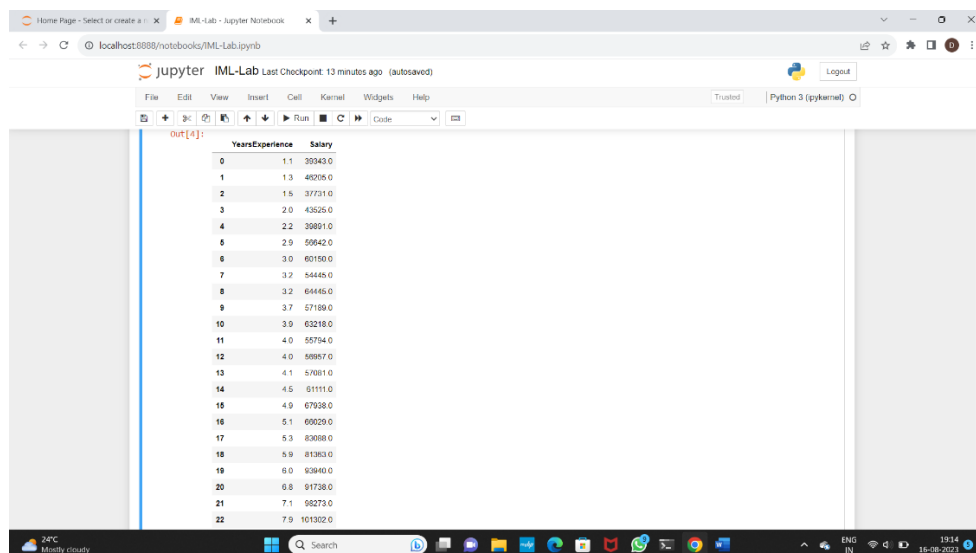
```
import matplotlib.pyplot as plt
```

```
import pandas as pd
```

```
dataset = pd.read_csv(r"C:\Users\dhana\Downloads\Salary_Data.csv")
```

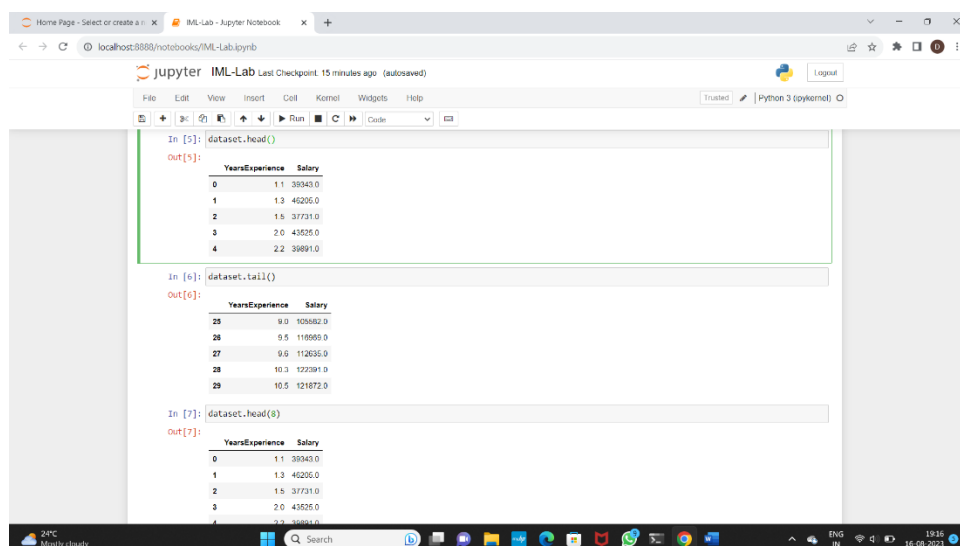
```
dataset
```

**Output:**



A screenshot of a Jupyter Notebook interface. The top bar shows 'jupyter IML-Lab' and 'Last Checkpoint: 13 minutes ago (autosaved)'. The notebook has a menu bar with 'File', 'Edit', 'View', 'Insert', 'Cell', 'Kernel', 'Widgets', and 'Help'. Below the menu is a toolbar with icons for file operations, running, and code execution. The main area shows a code cell with the command `dataset[4:]` and its output. The output is a table with two columns: 'YearsExperience' and 'Salary'. It displays rows 4 through 22 of the dataset.

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	50642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55754.0
12	4.0	59957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67928.0
16	5.1	60029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0



A screenshot of a Jupyter Notebook interface showing three code cells and their outputs. The first cell runs `dataset.head()` and shows the first 5 rows of the dataset. The second cell runs `dataset.tail()` and shows the last 5 rows. The third cell runs `dataset.head(3)` and shows the first 3 rows. The notebook interface is the same as the previous screenshot.

```
In [5]: dataset.head()
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

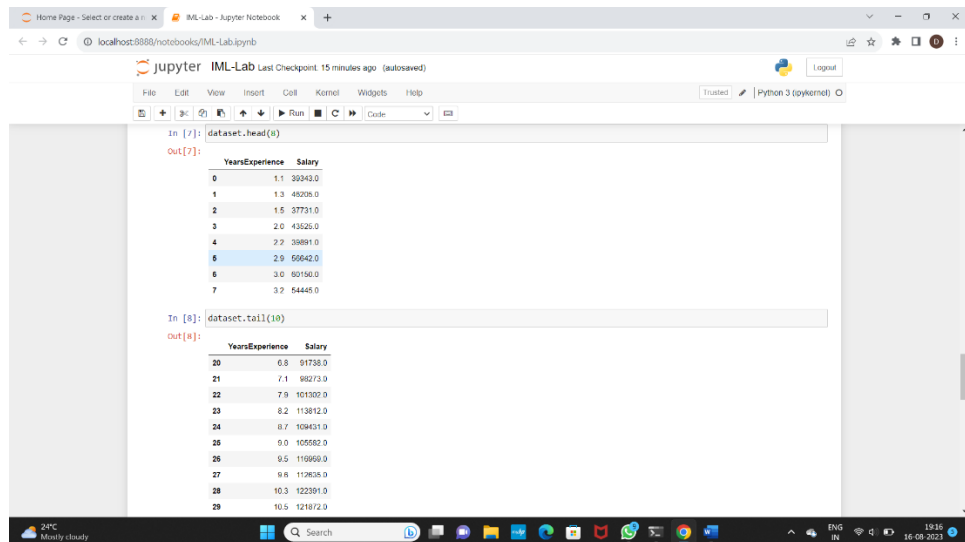
```
In [6]: dataset.tail()
```

	YearsExperience	Salary
25	9.0	103682.0
26	9.5	110569.0
27	9.6	113036.0
28	10.3	123981.0
29	10.5	121872.0

```
In [7]: dataset.head(3)
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0





```
import numpy as np
```

```
import matplotlib as mpl
```

```
from mpl_toolkits.mplot3d import Axes3D
```

```
import matplotlib.pyplot as plt
```

```
def generate_dataset(n):
```

```
    x = []
```

```
    y = []
```

```
    random_x1 = np.random.rand()
```

```
    random_x2 = np.random.rand()
```

```
    for i in range(n):
```

```
        x1 = i
```

```
        x2 = i/2 + np.random.rand()*n
```

```
        x.append([1, x1, x2])
```

```
        y.append(random_x1 * x1 + random_x2 * x2 + 1)
```

```
    return np.array(x), np.array(y)
```

```
x, y = generate_dataset(200)
```

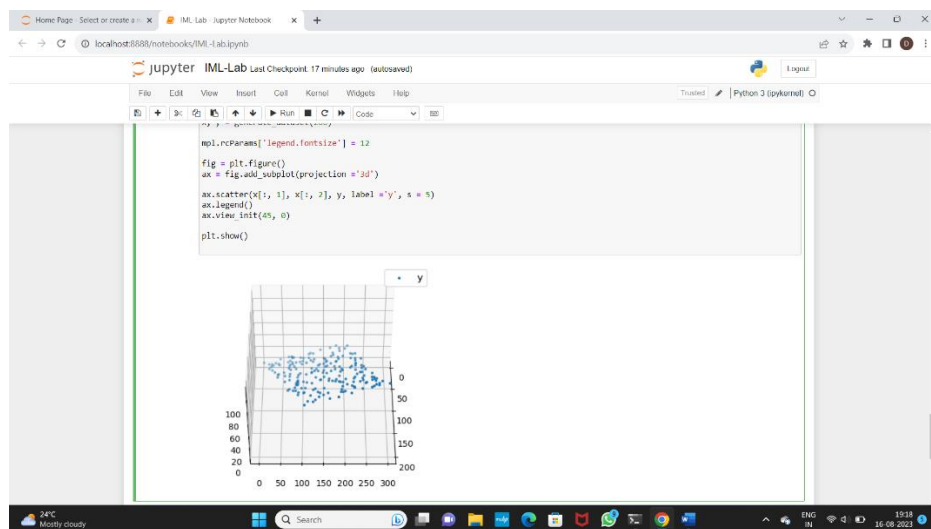
```
mpl.rcParams['legend.fontsize'] = 12
```

```
fig = plt.figure()
ax = fig.add_subplot(projection='3d')

ax.scatter(x[:, 1], x[:, 2], y, label='y', s = 5)
ax.legend()
ax.view_init(45, 0)

plt.show()
```

## Output:



```
from sklearn.linear_model import LinearRegression
import numpy as np
X = np.array([[1, 2, 3], [2, 3, 4], [3, 4, 5], [4, 5, 6]])
y = np.array([1, 2, 3, 4])
reg = LinearRegression()
reg.fit(X, y)
print(reg.coef_)
```

## Output:

```
[0.33333333 0.33333333 0.33333333]
```

# Logistic Regression

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

## Importing the dataset

```
In [4]: dataset = pd.read_csv(r'''C:\Users\dhana\OneDrive\Documents\Social_Network_Ads (
x = dataset.iloc[:, [2, 3]].values
y = dataset.iloc[:, -1].values
```

```
In [5]: dataset.head()
```

```
Out[5]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
In [6]: dataset.tail()
```

```
Out[6]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

## Splitting the dataset into the Training set and Test set

```
In [7]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train, y_test = train_test_split(x,y,test_size = 0.25,random_st
```

## Feature Scaling

```
In [8]: pip install sklearn
```

```
Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: sklearn in c:\users\dhana\appdata\roaming\python\python310\site-packages (0.0.post7)
Note: you may need to restart the kernel to use updated packages.
```

```
In [9]: pip install StandardScaler
```

```
Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: StandardScaler in c:\users\dhana\appdata\roaming\python\python310\site-packages (0.5)
Requirement already satisfied: scikit-elm in c:\users\dhana\appdata\roaming\python\python310\site-packages (from StandardScaler) (0.21a0)
Requirement already satisfied: pandas in c:\programdata\anaconda3\lib\site-packages (from StandardScaler) (1.5.3)
Requirement already satisfied: scikit-learn in c:\programdata\anaconda3\lib\site-packages (from StandardScaler) (1.2.1)
Requirement already satisfied: numpy in c:\programdata\anaconda3\lib\site-packages (from StandardScaler) (1.23.5)
Requirement already satisfied: dask in c:\programdata\anaconda3\lib\site-packages (from StandardScaler) (2022.7.0)
Requirement already satisfied: partd>=0.3.10 in c:\programdata\anaconda3\lib\site-packages (from dask->StandardScaler) (1.2.0)
Requirement already satisfied: packaging>=20.0 in c:\programdata\anaconda3\lib\site-packages (from dask->StandardScaler) (22.0)
Requirement already satisfied: fsspec>=0.6.0 in c:\programdata\anaconda3\lib\site-packages (from dask->StandardScaler) (2022.11.0)
Requirement already satisfied: cloudpickle>=1.1.1 in c:\programdata\anaconda3\lib\site-packages (from dask->StandardScaler) (2.0.0)
Requirement already satisfied: toolz>=0.8.2 in c:\programdata\anaconda3\lib\site-packages (from dask->StandardScaler) (0.12.0)
Requirement already satisfied: pyyaml>=5.3.1 in c:\programdata\anaconda3\lib\site-packages (from dask->StandardScaler) (6.0)
Requirement already satisfied: pytz>=2020.1 in c:\programdata\anaconda3\lib\site-packages (from pandas->StandardScaler) (2022.7)
Requirement already satisfied: python-dateutil>=2.8.1 in c:\programdata\anaconda3\lib\site-packages (from pandas->StandardScaler) (2.8.2)
Requirement already satisfied: scipy in c:\programdata\anaconda3\lib\site-packages (from scikit-elm->StandardScaler) (1.10.0)
Requirement already satisfied: joblib>=1.1.1 in c:\programdata\anaconda3\lib\site-packages (from scikit-learn->StandardScaler) (1.1.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\programdata\anaconda3\lib\site-packages (from scikit-learn->StandardScaler) (2.2.0)
Requirement already satisfied: locket in c:\programdata\anaconda3\lib\site-packages (from partd>=0.3.10->dask->StandardScaler) (1.0.0)
Requirement already satisfied: six>=1.5 in c:\programdata\anaconda3\lib\site-packages (from python-dateutil>=2.8.1->pandas->StandardScaler) (1.16.0)
Note: you may need to restart the kernel to use updated packages.
```

```
In [10]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

## Training the Logistic Regression model on training set

```
In [11]: from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression(random_state = 0)
classifier.fit(X_train,y_train)
```

```
Out[11]: LogisticRegression
LogisticRegression(random_state=0)
```

## Predicting the Test set results

```
In [12]: y_pred = classifier.predict(X_test)
```

## Making the confusion Matrix


```
In [13]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test,y_pred)
print(cm)
```

```
[[65  3]
 [ 8 24]]
```

```
In [17]: from matplotlib.colors import ListedColormap
X_set, y_set = X_train, y_train
X1,X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() -1, stop = X_set[:, 0].max() +1, step = 0.5),
                    np.arange(start = X_set[:, 1].min() -1, stop = X_set[:, 1].max() +1, step = 0.5))
plt.contourf(X1,X2,classifier.predict(np.array([X1.ravel(),X2.ravel()]).T).reshape(X1.shape),
             alpha = 0.75,cmap = ListedColormap(('red','green')))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j,0],X_set[y_set == j,1],
                c = ListedColormap(('red','green'))(i),label = j)
plt.title('Logistic Regression (Training set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

C:\Users\dhana\AppData\Local\Temp\ipykernel\_21176\3711601667.py:10: UserWarning:  
\*c\* argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with \*x\* & \*y\*. Please use the \*color\* keyword-argument or provide a 2D array with a single row if you intend to specify the same RGB or RGBA value for all points.

```
plt.scatter(X_set[y_set == j,0],X_set[y_set == j,1],
```



In [ ]:

# Multiclass Classification using Decision Trees

## Importing Required Libraries

```
In [1]: import numpy as np
        from sklearn.datasets import load_iris
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import confusion_matrix
```

## Dataset

```
In [2]: Iris_dataset = load_iris()
```

## Training and Testing

```
In [3]: X_train, X_test, y_train, y_test = train_test_split(Iris_dataset.data, Iris_dataset.target, rand
```

## Training using Decision Tree Classifier

```
In [4]: clf = DecisionTreeClassifier()
        clf.fit(X_train, y_train)
```

```
Out[4]: DecisionTreeClassifier()
```

## Prediction

```
In [5]: y_pred = clf.predict(X_test)
```

## Accuracy

```
In [6]: accuracy = np.mean(y_pred == y_test)
        print("Accuracy is {:.2f} %".format(accuracy*100))
```

Accuracy is 97.37 %

# K-Nearest Neighbor(KNN)

importing libraries and dataset

```
import numpy as nm
import matplotlib.pyplot as mtp
import pandas as pd

data_set= pd.read_csv(r'''C:\Users\katka\Downloads\
Social_Network_Ads.csv''')
```

Extracting Independent and dependent Variable

```
x= data_set.iloc[:, [2,3]].values
y= data_set.iloc[:, 4].values
```

Splitting the dataset into training and test set.

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test= train_test_split(x, y, test_size=
0.25, random_state=0)
```

feature Scaling

```
from sklearn.preprocessing import StandardScaler
st_x= StandardScaler()
x_train= st_x.fit_transform(x_train)
x_test= st_x.transform(x_test)
```

Fitting K-NN classifier to the training set

```
from sklearn.neighbors import KNeighborsClassifier
classifier= KNeighborsClassifier(n_neighbors=5, metric='minkowski',
p=2 )
classifier.fit(x_train, y_train)

KNeighborsClassifier()
```

Predicting the test set result

```
y_pred= classifier.predict(x_test)
```

Creating the Confusion matrix

```
from sklearn.metrics import confusion_matrix
cm= confusion_matrix(y_test, y_pred)
```




## Visualizing the training set result

```
from matplotlib.colors import ListedColormap
x_set, y_set = x_train, y_train
x1, x2 = nm.meshgrid(nm.arange(start = x_set[:, 0].min() - 1, stop =
x_set[:, 0].max() + 1, step = 0.01),
nm.arange(start = x_set[:, 1].min() - 1, stop = x_set[:, 1].max() + 1,
step = 0.01))
mtp.contourf(x1, x2, classifier.predict(nm.array([x1.ravel(),
x2.ravel()]).T).reshape(x1.shape),
alpha = 0.75, cmap = ListedColormap(('red', 'green' )))
mtp.xlim(x1.min(), x1.max())
mtp.ylim(x2.min(), x2.max())
for i, j in enumerate(nm.unique(y_set)):
    mtp.scatter(x_set[y_set == j, 0], x_set[y_set == j, 1],
               c = ListedColormap(('red', 'green'))(i), label = j)
mtp.title('K-NN Algorithm (Training set)')
mtp.xlabel('Age')
mtp.ylabel('Estimated Salary')
mtp.legend()
mtp.show()
```

*\*c\** argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with *\*x\** & *\*y\**. Please use the *\*color\** keyword-argument or provide a 2-D array with a single row if you intend to specify the same RGB or RGBA value for all points.

*\*c\** argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with *\*x\** & *\*y\**. Please use the *\*color\** keyword-argument or provide a 2-D array with a single row if you intend to specify the same RGB or RGBA value for all points.



# Naive Bayes

## Importing the libraries

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

## Importing the dataset

```
In [2]: dataset = pd.read_csv('Social_Network_Ads.csv')
X = dataset.iloc[:, [2, 3]].values
v = dataset.iloc[:, -1].values
```

```
In [12]: print(dataset)
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
..	...	...	...	...	...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

[400 rows x 5 columns]

-----  
-  
**AttributeError** Traceback (most recent call last)

Cell In[12], line 1  
----> 1 print(dataset).head(50)

**AttributeError**: 'NoneType' object has no attribute 'head'

## Splitting the dataset into the Training set and Test set

```
In [3]: from sklearn.model_selection import train_test_split
X_train, X_test, v_train, v_test = train_test_split(X, v, test_size = 0.25)
```

## Feature Scaling

```
In [4]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

## ## Training the Naive Bayes model on the Training set

```
In [5]: from sklearn.naive_bayes import GaussianNB
classifier = GaussianNB()
classifier.fit(X_train, y_train)
```

```
Out[5]:
```

▼ GaussianNB
GaussianNB()

## Predicting the Test set results

```
In [6]: y_pred = classifier.predict(X_test)
```

## Making the Confusion Matrix

```
In [7]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
print(cm)
```


```
[[65  3]
 [ 7 25]]
```

## Visualising the Training set results

```
In [8]: from matplotlib.colors import ListedColormap
X_set, y_set = X_train, y_train
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[
                                np.arange(start = X_set[:, 1].min() - 1, stop = X_set[
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()])).
                alpha = 0.75, cmap = ListedColormap(['red', 'green']))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(['red', 'green'])(i), label = j)
plt.title('Naive Bayes (Training set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

C:\Users\dyp\AppData\Local\Temp\ipykernel\_4772\2643083737.py:10: UserWarning: \*c\* argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with \*x\* & \*y\*. Please use the \*color\* keyword-argument or provide a 2D array with a single row if you intend to specify the same RGB or RGBA value for all points.

```
plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
```




**Visualising the Test set results**

```
In [9]: from matplotlib.colors import ListedColormap
X_set, y_set = X_test, y_test
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[
                                np.arange(start = X_set[:, 1].min() - 1, stop = X_set[
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()])).
                alpha = 0.75, cmap = ListedColormap(['red', 'green']))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(['red', 'green'])(i), label = j)
plt.title('Naive Bayes (Test set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

C:\Users\dyp\AppData\Local\Temp\ipykernel\_4772\664088336.py:10: UserWarning: \*c\* argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with \*x\* & \*y\*. Please use the \*color\* keyword-argument or provide a 2D array with a single row if you intend to specify the same RGB or RGBA value for all points.

```
plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
```



# Support Vector Machine (SVM)

## Importing the libraries

```
In [0]: ▶ import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

## Importing the dataset

```
In [0]: ▶ dataset = pd.read_csv('Social_Network_Ads.csv')
X = dataset.iloc[:, [2, 3]].values
y = dataset.iloc[:, -1].values
```

## Splitting the dataset into the Training set and Test set

```
In [0]: ▶ from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.
```

## Feature Scaling

```
In [0]: ▶ from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

## Training the SVM model on the Training set

```
In [5]: ▶ from sklearn.svm import SVC
classifier = SVC(kernel = 'linear', random_state = 0)
classifier.fit(X_train, y_train)
```

```
Out[5]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=
0.0,
        decision_function_shape='ovr', degree=3, gamma='scale', kernel='lin
ear',
        max_iter=-1, probability=False, random_state=0, shrinking=True, tol
=0.001,
        verbose=False)
```

## Predicting the Test set results

```
In [0]: ► y_pred = classifier.predict(X_test)
```

## Making the Confusion Matrix

```
In [7]: ► from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_pred)  
print(cm)
```

```
[[66  2]  
 [ 8 24]]
```




## Visualising the Training set results

```
In [8]: ▶ from matplotlib.colors import ListedColormap
X_set, y_set = X_train, y_train
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_s
                        np.arange(start = X_set[:, 1].min() - 1, stop = X_s
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()
                        alpha = 0.75, cmap = ListedColormap(('red', 'green'))))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(('red', 'green'))(i), label = j)
plt.title('SVM (Training set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.




## Visualising the Test set results

```
In [9]: ▶ from matplotlib.colors import ListedColormap
X_set, y_set = X_test, y_test
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_s
                        np.arange(start = X_set[:, 1].min() - 1, stop = X_s
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()
                        alpha = 0.75, cmap = ListedColormap(('red', 'green'))))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(('red', 'green'))(i), label = j)
plt.title('SVM (Test set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.



# Decision Tree Classification

## Importing the libraries

```
In [0]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

## Importing the dataset

```
In [0]: dataset = pd.read_csv('Social_Network_Ads.csv')
X = dataset.iloc[:, [2, 3]].values
y = dataset.iloc[:, -1].values
```

## Splitting the dataset into the Training set and Test set

```
In [0]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25,
```

## Feature Scaling

```
In [0]: from sklearn.preprocessing import StandardScaler
        sc = StandardScaler()
        X_train = sc.fit_transform(X_train)
        X_test = sc.transform(X_test)
```

## Training the Decision Tree Classification model on the Training set

```
In [5]: from sklearn.tree import DecisionTreeClassifier
classifier = DecisionTreeClassifier(criterion = 'entropy', random_state = 0)
classifier.fit(X_train, y_train)
```

[illegible]

## Predicting the Test set results

```
In [0]: ► y_pred = classifier.predict(X_test)
```

## Making the Confusion Matrix

```
In [7]: ► from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_pred)  
print(cm)
```


```
[[62  6]  
 [ 3 29]]
```

## Visualising the Training set results

```
In [8]: ▶ from matplotlib.colors import ListedColormap
X_set, y_set = X_train, y_train
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.5),
                     np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.5))
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()])).reshape(X1.shape),
             alpha = 0.75, cmap = ListedColormap(('red', 'green')))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(('red', 'green'))(i), label = j)
plt.title('Decision Tree Classification (Training set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.




## Visualising the Test set results

```
In [9]: ▶ from matplotlib.colors import ListedColormap
X_set, y_set = X_test, y_test
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.5),
                     np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.5))
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()])).
             .reshape(X1.shape),
             alpha = 0.75, cmap = ListedColormap(('red', 'green')))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(('red', 'green'))(i), label = j)
plt.title('Decision Tree Classification (Test set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.



# Random Forest Classification

## Importing the libraries

```
In [0]: ▶ import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

## Importing the dataset

```
In [0]: ▶ dataset = pd.read_csv('Social_Network_Ads.csv')
X = dataset.iloc[:, [2, 3]].values
y = dataset.iloc[:, -1].values
```

## Splitting the dataset into the Training set and Test set

```
In [0]: ▶ from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.
```

## Feature Scaling

```
In [0]: ▶ from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

## Training the Random Forest Classification model on the Training set

```
In [5]: ▶ from sklearn.ensemble import RandomForestClassifier
classifier = RandomForestClassifier(n_estimators = 10, criterion = 'entropy')
classifier.fit(X_train, y_train)
```

```
Out[5]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                               criterion='entropy', max_depth=None, max_features='auto',
                               max_leaf_nodes=None, max_samples=None,
                               min_impurity_decrease=0.0, min_impurity_split=None,
                               min_samples_leaf=1, min_samples_split=2,
                               min_weight_fraction_leaf=0.0, n_estimators=10,
                               n_jobs=None, oob_score=False, random_state=0, verbose=0,
                               warm_start=False)
```

## Predicting the Test set results

```
In [0]: ► y_pred = classifier.predict(X_test)
```

## Making the Confusion Matrix

```
In [7]: ► from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_pred)  
print(cm)
```

```
[[63  5]  
 [ 4 28]]
```




## Visualising the Training set results

```
In [8]: ▶ from matplotlib.colors import ListedColormap
X_set, y_set = X_train, y_train
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_s
                        np.arange(start = X_set[:, 1].min() - 1, stop = X_s
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()
                        alpha = 0.75, cmap = ListedColormap(('red', 'green'))))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(('red', 'green'))(i), label = j)
plt.title('Random Forest Classification (Training set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.



## Visualising the Test set results

```
In [9]: ▶ from matplotlib.colors import ListedColormap
X_set, y_set = X_test, y_test
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_s
                        np.arange(start = X_set[:, 1].min() - 1, stop = X_s
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()
                        alpha = 0.75, cmap = ListedColormap(('red', 'green'))))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(('red', 'green'))(i), label = j)
plt.title('Random Forest Classification (Test set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.

